

# Rain Prediction Using ANN ( Tensorflow 2 - Keras )

```
In [34]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df = pd.read_csv("D:\data science\Data scientist\Projects\P5 Rain pred using ANN\weatherAUS.csv")
```

## Data Cleaning

```
In [3]: df.head()
```

```
Out[3]:
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	WindGustSpeed	WindDir9am	...	Humidity9am	Humidity3pm	Pre
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	W	44.0	W	...	71.0	22.0	
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	WNW	44.0	NNW	...	44.0	25.0	
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	WSW	46.0	W	...	38.0	30.0	
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	NE	24.0	SE	...	45.0	16.0	
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	W	41.0	ENE	...	82.0	33.0	

5 rows × 23 columns



```
In [4]: len(df)
```

```
Out[4]: 145460
```

```
In [5]: df.head().isnull().sum()
```

```
Out[5]: Date                0
Location                  0
MinTemp                   0
MaxTemp                   0
Rainfall                  0
Evaporation                5
Sunshine                  5
WindGustDir               0
WindGustSpeed             0
WindDir9am                0
WindDir3pm               0
WindSpeed9am              0
WindSpeed3pm              0
Humidity9am               0
Humidity3pm               0
Pressure9am               0
Pressure3pm               0
Cloud9am                  3
Cloud3pm                  3
Temp9am                   0
Temp3pm                   0
RainToday                 0
RainTomorrow              0
dtype: int64
```

```
In [6]: df = df.dropna()
```

```
In [7]: df.head().isnull().sum()
```

```
Out[7]: Date          0
Location          0
MinTemp          0
MaxTemp          0
Rainfall         0
Evaporation       0
Sunshine         0
WindGustDir       0
WindGustSpeed     0
WindDir9am       0
WindDir3pm       0
WindSpeed9am     0
WindSpeed3pm     0
Humidity9am      0
Humidity3pm      0
Pressure9am      0
Pressure3pm      0
Cloud9am         0
Cloud3pm         0
Temp9am          0
Temp3pm          0
RainToday        0
RainTomorrow     0
dtype: int64
```

```
In [8]: df.head()
```

Out[8]:

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	WindGustSpeed	WindDir9am	...	Humidity9am	Humidity3pm
6049	2009-01-01	Cobar	17.9	35.2	0.0	12.0	12.3	SSW	48.0	ENE	...	20.0	13.0
6050	2009-01-02	Cobar	18.4	28.9	0.0	14.8	13.0	S	37.0	SSE	...	30.0	8.0
6052	2009-01-04	Cobar	19.4	37.6	0.0	10.8	10.6	NNE	46.0	NNE	...	42.0	22.0
6053	2009-01-05	Cobar	21.9	38.4	0.0	11.4	12.2	WNW	31.0	WNW	...	37.0	22.0
6054	2009-01-06	Cobar	24.2	41.0	0.0	11.2	8.4	WNW	35.0	NW	...	19.0	15.0


5 rows × 23 columns

```
In [9]: df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
Index: 56420 entries, 6049 to 142302
Data columns (total 23 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   Date                56420 non-null object
1   Location            56420 non-null object
2   MinTemp             56420 non-null float64
3   MaxTemp             56420 non-null float64
4   Rainfall            56420 non-null float64
5   Evaporation         56420 non-null float64
6   Sunshine            56420 non-null float64
7   WindGustDir         56420 non-null object
8   WindGustSpeed       56420 non-null float64
9   WindDir9am          56420 non-null object
10  WindDir3pm          56420 non-null object
11  WindSpeed9am        56420 non-null float64
12  WindSpeed3pm        56420 non-null float64
13  Humidity9am         56420 non-null float64
14  Humidity3pm         56420 non-null float64
15  Pressure9am         56420 non-null float64
16  Pressure3pm         56420 non-null float64
17  Cloud9am            56420 non-null float64
18  Cloud3pm            56420 non-null float64
19  Temp9am             56420 non-null float64
20  Temp3pm             56420 non-null float64
21  RainToday           56420 non-null object
22  RainTomorrow        56420 non-null object
dtypes: float64(16), object(7)
memory usage: 10.3+ MB
```

## Data Visualization


```
In [52]: plt.figure(figsize=(12,6),dpi = 150)
sns.countplot(data = df,x='RainToday')
plt.tight_layout()
```



```
In [81]: plt.figure(figsize=(12,12),dpi = 150)
sns.heatmap(df.corr(),annot = True,fmt='.2g',annot_kws={"size": 6})
plt.tight_layout()
```



```
In [80]: lowest = df.groupby(by='Location')['MinTemp'].min().reset_index()
plt.figure(figsize=(8, 6), dpi = 150)
sns.barplot(data=lowest, x='Location', y='MinTemp')
plt.tight_layout()
```



## Data Preparation

```
In [10]: object_columns = df.select_dtypes(include=['object']).columns.tolist()
object_columns
```

```
Out[10]: ['Date',
'Location',
'WindGustDir',
'WindDir9am',
'WindDir3pm',
'RainToday',
'RainTomorrow']
```

```
In [11]: from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
for i in object_columns:
    df[i] = label_encoder.fit_transform(df[i])

df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 56420 entries, 6049 to 142302
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  56420 non-null  int32
1   Location              56420 non-null  int32
2   MinTemp               56420 non-null  float64
3   MaxTemp               56420 non-null  float64
4   Rainfall              56420 non-null  float64
5   Evaporation           56420 non-null  float64
6   Sunshine              56420 non-null  float64
7   WindGustDir           56420 non-null  int32
8   WindGustSpeed         56420 non-null  float64
9   WindDir9am            56420 non-null  int32
10  WindDir3pm            56420 non-null  int32
11  WindSpeed9am          56420 non-null  float64
12  WindSpeed3pm          56420 non-null  float64
13  Humidity9am           56420 non-null  float64
14  Humidity3pm           56420 non-null  float64
15  Pressure9am           56420 non-null  float64
16  Pressure3pm           56420 non-null  float64
17  Cloud9am              56420 non-null  float64
18  Cloud3pm              56420 non-null  float64
19  Temp9am               56420 non-null  float64
20  Temp3pm               56420 non-null  float64
21  RainToday             56420 non-null  int32
22  RainTomorrow          56420 non-null  int32
dtypes: float64(16), int32(7)
memory usage: 8.8 MB
```

```
In [12]: X = df.drop(['RainTomorrow'],axis = 1)
y = df['RainTomorrow']
```

```
In [13]: X.head(7)
```

```
Out[13]:
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	WindGustSpeed	WindDir9am	...	WindSpeed3pm	Humidity9am
6049	407	4	17.9	35.2	0.0	12.0	12.3	11	48.0	1	...	20.0	20.0
6050	408	4	18.4	28.9	0.0	14.8	13.0	8	37.0	10	...	19.0	30.0
6052	410	4	19.4	37.6	0.0	10.8	10.6	5	46.0	5	...	15.0	42.0
6053	411	4	21.9	38.4	0.0	11.4	12.2	14	31.0	14	...	6.0	37.0
6054	412	4	24.2	41.0	0.0	11.2	8.4	14	35.0	7	...	13.0	19.0
6055	413	4	27.1	36.1	0.0	13.0	0.0	3	43.0	3	...	20.0	26.0
6056	414	4	23.3	34.0	0.0	9.8	12.6	11	41.0	8	...	19.0	33.0

7 rows × 22 columns



```
In [14]: print(y.unique())
```

```
[0 1]
```

```
In [15]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)

X.shape
```

```
Out[15]: (56420, 22)
```

```
In [16]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaled_X_train = scaler.fit_transform(X_train)
scaled_X_test = scaler.transform(X_test)
```

```
In [17]: print(len(y_train))
         print(len(scaled_X_train))
         print(len(scaled_X_test))
         print(len(y_test))
```

```
45136
45136
11284
11284
```

## Model Building

```
In [29]: from keras.layers import Dense, BatchNormalization, Dropout
from keras.models import Sequential
from keras import callbacks
from keras.optimizers import Adam

early_stopping = callbacks.EarlyStopping(
    min_delta=0.001,
    patience=10,
    restore_best_weights=True,
)

model = Sequential()

model.add(Dense(units = 32, kernel_initializer = 'uniform', activation = 'relu', input_dim = 22))
model.add(Dense(units = 32, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dense(units = 16, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dropout(0.25))
model.add(Dense(units = 16, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dropout(0.25))
model.add(Dense(units = 8, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dropout(0.5))
model.add(Dense(units = 8, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dropout(0.5))
model.add(Dense(units = 1, kernel_initializer = 'uniform', activation = 'sigmoid'))

opt = Adam(learning_rate=0.001)
model.compile(optimizer = opt, loss = 'binary_crossentropy', metrics = ['accuracy'])

history = model.fit(scaled_X_train,y_train, batch_size = 32, epochs = 150, callbacks=[early_stopping], validation_split=0.2
```



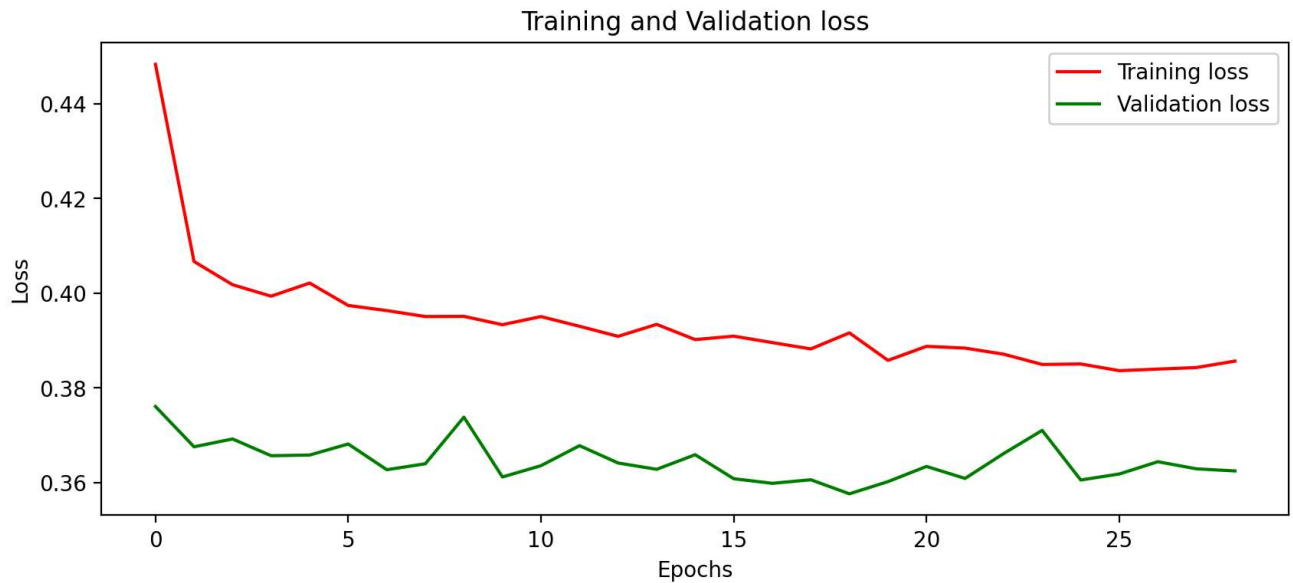
Epoch 1/150  
1129/1129 [=====] - 3s 2ms/step - loss: 0.4483 - accuracy: 0.7804 - val\_loss: 0.3761 - val\_accuracy: 0.7771  
Epoch 2/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.4067 - accuracy: 0.7804 - val\_loss: 0.3676 - val\_accuracy: 0.7771  
Epoch 3/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.4018 - accuracy: 0.7804 - val\_loss: 0.3692 - val\_accuracy: 0.7771  
Epoch 4/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3994 - accuracy: 0.7804 - val\_loss: 0.3657 - val\_accuracy: 0.7771  
Epoch 5/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.4021 - accuracy: 0.7804 - val\_loss: 0.3658 - val\_accuracy: 0.7771  
Epoch 6/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3974 - accuracy: 0.7804 - val\_loss: 0.3682 - val\_accuracy: 0.7771  
Epoch 7/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3963 - accuracy: 0.7804 - val\_loss: 0.3627 - val\_accuracy: 0.7771  
Epoch 8/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3951 - accuracy: 0.7926 - val\_loss: 0.3640 - val\_accuracy: 0.8501  
Epoch 9/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3951 - accuracy: 0.8186 - val\_loss: 0.3738 - val\_accuracy: 0.7771  
Epoch 10/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3933 - accuracy: 0.8156 - val\_loss: 0.3612 - val\_accuracy: 0.8530  
Epoch 11/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3951 - accuracy: 0.8293 - val\_loss: 0.3636 - val\_accuracy: 0.8490  
Epoch 12/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3930 - accuracy: 0.8327 - val\_loss: 0.3678 - val\_accuracy: 0.8475  
Epoch 13/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3909 - accuracy: 0.8339 - val\_loss: 0.3641 - val\_accuracy: 0.8531  
Epoch 14/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3934 - accuracy: 0.8406 - val\_loss: 0.3628 - val\_accuracy: 0.8508  
Epoch 15/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3902 - accuracy: 0.8402 - val\_loss: 0.3659 - val\_accuracy: 0.8463  
Epoch 16/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3909 - accuracy: 0.8387 - val\_loss: 0.3608 - val\_accuracy: 0.8532  
Epoch 17/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3896 - accuracy: 0.8400 - val\_loss: 0.3598 - val\_accuracy: 0.8530  
Epoch 18/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3882 - accuracy: 0.8397 - val\_loss: 0.3606 - val\_accuracy: 0.8511  
Epoch 19/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3916 - accuracy: 0.8401 - val\_loss: 0.3576 - val\_accuracy: 0.8535  
Epoch 20/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3858 - accuracy: 0.8420 - val\_loss: 0.3602 - val\_accuracy: 0.8542  
Epoch 21/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3888 - accuracy: 0.8413 - val\_loss: 0.3634 - val\_accuracy: 0.8533  
Epoch 22/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3884 - accuracy: 0.8391 - val\_loss: 0.3609 - val\_accuracy: 0.8517  
Epoch 23/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3871 - accuracy: 0.8412 - val\_loss: 0.3661 - val\_accuracy: 0.8485  
Epoch 24/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3849 - accuracy: 0.8424 - val\_loss: 0.3710 - val\_accuracy: 0.8506  
Epoch 25/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3851 - accuracy: 0.8396 - val\_loss: 0.3605 - val\_accuracy: 0.8501  
Epoch 26/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3836 - accuracy: 0.8412 - val\_loss: 0.3618 - val\_accuracy: 0.8488  
Epoch 27/150  
1129/1129 [=====] - 2s 2ms/step - loss: 0.3840 - accuracy: 0.8385 - val\_loss: 0.3644 - val\_accuracy: 0.8405  
Epoch 28/150

```
1129/1129 [=====] - 2s 2ms/step - loss: 0.3843 - accuracy: 0.8381 - val_loss: 0.3629 - val_accuracy: 0.8497
Epoch 29/150
1129/1129 [=====] - 2s 2ms/step - loss: 0.3857 - accuracy: 0.8424 - val_loss: 0.3625 - val_accuracy: 0.8502
```

```
In [30]: history_df = pd.DataFrame(history.history)
plt.figure(figsize=(10,4),dpi=200)
plt.plot(history_df.loc[:, ['loss']], "Red", label='Training loss')
plt.plot(history_df.loc[:, ['val_loss']], "Green", label='Validation loss')
plt.title('Training and Validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc="best")

plt.show()

print(len(y_train))
print(len(scaled_X_train))
print(len(scaled_X_test))
print(len(y_test))
```



```
45136
45136
11284
11284
```


## Model Prediction

```
In [31]: y_pred = model.predict(scaled_X_test)
y_pred = (y_pred > 0.5)
```

```
353/353 [=====] - 0s 930us/step
```

```
In [32]: import seaborn as sns
from sklearn.metrics import confusion_matrix, classification_report
plt.subplots(figsize=(12,8))
cf_matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(cf_matrix/np.sum(cf_matrix), annot = True, annot_kws = {'size':15})
```

Out[32]: <Axes: >



```
In [33]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.89	0.92	0.91	8799
1	0.69	0.61	0.65	2485
accuracy			0.85	11284
macro avg	0.79	0.77	0.78	11284
weighted avg	0.85	0.85	0.85	11284

## Results

**Model accuracy = 85%**

**No rain prediction f1-score = 0.91**

**Will rain prediction f1-score = 0.65**